

Assignment: System Integrity Verifier (SIV)

Introduction

In this assignment you will implement a very simple system integrity verifier (SIV) for a Linux system. The goal of the SIV is to detect file system modifications occurring within a directory tree. The SIV outputs statistics and warnings about changes to a report file specified by the user.

The SIV can be run either in *initialization mode* or in *verification mode*.

Initialization mode

In initialization mode, the SIV program requires the user to enter a path to a monitored directory, another path to a verification file (outside the monitored directory), a third path to a report file and a hash function (e.g., the SIV should support at least MD-5 and SHA-1 message digests). The program will do the following:

- a) Verify that the specified monitored directory exists
- b) Verify that the specified verification file and the report file are outside the monitored directory
- c) Verify that the specified hash function is supported by your SIV
- d) If the verification file or report file exists already then the user will be asked if he wants to overwrite the existing file. If answer is “no”, the program terminates.
- e) Recursively iterate through the directory contents (of arbitrary depth level) and for each file found record the following information in the verification file (it is up to you decide the format used to store records in the verification file):
 - I. Full path to file, including filename
 - II. Size of the file
 - III. Name of user owning the file
 - IV. Name of group owning the file
 - V. Access rights to the file (either octal or symbolic)
 - VI. Last modification date
 - VII. Computed message digest with specified hash function over file contents
- f) Write to the report file (this must be a text file) a summary of your findings:
 - I. Full pathname to monitored directory
 - II. Full pathname to verification file
 - III. Number of directories parsed
 - IV. Number of files parsed
 - V. Time to complete the initialization mode

Verification mode

In verification mode the user provides the path of the verification file and a path to a report file. The program will do the following:

- a) Verify that the specified verification file exists and, if true, begin parsing the file

- b) Verify that the specified verification file is outside the monitored directory
- c) Verify that the specified report file is outside the monitored directory
- d) Recursively iterate through the directory contents (of arbitrary depth level) and verify that it matches exactly the information from the verification file. A warning must be printed to the report file for each entry that diverges from the records in the verification file, along with information about what is different
 - I. New or removed files/directories
 - II. Files with a different size than recorded
 - III. Files with a different message digest than computed before
 - IV. Files/directories with a different user/group
 - V. Files/directories with modified access right
 - VI. Files/directories with a different modification date
- e) Write to the report file (this must be a text file) a summary of your findings
 - I. Full pathname to monitored directory
 - II. Full pathname to verification file
 - III. Full pathname to report file
 - IV. Number of directories parsed
 - V. Number of files parsed
 - VI. Number of warning issued
 - VII. Time to complete the verification mode

User interface

The program must be runnable from the command-line and accept the following command-line arguments:

```
siv <-i|-v|-h> -D <monitored_directory> -V <verification_file> -R <report_file> -H
<hash_function>
```

The options -i (indicating initialization mode), -v (indicating verification mode), and -h (indicating help mode) are mutually exclusive.

The options -V and -H are mutually exclusive, meaning that you specify the hash function only when you create the verification file (in initialization mode). In verification mode, the hash function must be recovered from the verification file.

When the help option (-h) is given, the program will print the accepted command-line arguments with a short explanation for each, and show an example how to run the program in initialization mode and verification mode respectively. In addition, all supported hash functions must be listed with the syntax required to specify them on the command-line.

No other command-line arguments than those describe in this section should be required for executing the SIV in initialization or verification mode. Please, don't invent your own set of command-line arguments.

Example 1: Initialization mode

```
siv -i -D important_directory -V verificationDB -R my_report.txt -H sha1
```

Example 2: Verification mode

```
siv -v -D important_directory -V verificationDB -R my_report2.txt
```

Implementation

You are free to use Bash shell scripting or one of the following programming languages to implement the SIV (or a combination thereof): C, C++, Objective-C, Java, Perl, or Python. You are not allowed to use libraries of functions that are not included in the standard installation of the programming language. Using an already existing SIV system (such as Tripwire) is not allowed either.

It is strongly suggested that you consult the man pages for stat(1) and stat(2) shown in references. Stat(1) is useful to use from a shell script and stat(2) is useful from within a C-program. Other programming languages may provide simplified interfaces to stat(2) or their own API.

If you have a Windows or MacOS computer you can install VirtualBox from <https://www.virtualbox.org/>, create a Linux VM and then install the Linux Ubuntu Desktop <http://www.ubuntu.com/> inside the VM.

Testing

The best way to test your SIV is to download the complete Linux kernel source from <http://www.kernel.org>. You can uncompress the source code XZ archive with the command:

```
tar xvfJ <filename.tar.xz>
```

You should try changing files as outlined above under section “Verification mode” and check that your SIV can find the changes.

Deliverables and evaluation

This is an individual assignment with two deliverables. You will deliver 1) a lab report (as PDF file) and 2) the complete SIV source code as a tar.gz archive. In addition, you will include in the tar.gz archive one sample verification file and one sample report created by your program showing detected changes as part of the testing above. The sample report and the sample verification must originate from the same verification run of the SIV and should demonstrate the ability to detect all type of changes described under section “Verification mode”.

The delivered source code will be compiled to an executable if necessary. It is therefore imperative that your code compiles (unless it is a script) and that you document which compiler/interpreter version must be used. Specific sets of data will be used to test the functionality of the code. You will not have access to these sets of data during code development.

Authors who submit very similar code will be asked to explain in detail their programs in order to eliminate plagiarism suspicions. It is expected that authors can explain every line in the submitted program as well as the format of the verification file.

The report will be written with Microsoft Word and formatted according to the ACM small standard format available from <https://www.acm.org/publications/acm-word-style-guide> (first template on that page). The report will be uploaded to Its Learning in PDF format (not Word).

The report must have the following structure:

1. Introduction.
Here you explain the goal of the assignment and the possible applications of the SIV
2. Design and implementation
Here you explain how you detect each type of change (all 6 types) that can occur to the records. I expect detailed algorithm descriptions with pseudo-code. Also, I expect you to specify in detail the format of the verification file. The level of detail must be such that one should be able to write a program that interprets the contents of the verification file without guessing or looking into your code. For example, if each record (i.e., file-related data) occupies one row in the verification file, describe the meaning of each field including the appropriate data type for it (e.g., string terminated by zero, integer, floating point, hex string etc.). Explain also which programming language you have chosen and why. If there are any software dependencies required to run your program, make sure you specify them so I can prepare a similar environment when testing your program.
3. Usage
A short manual how to use your SIV, including examples for the most common operations,
4. Limitations
Explain if your program has any limitations, if any, in terms of the tasks it must perform. These limitations must be limited special cases (e.g., special characters in file names) as your program must implement all functionality described in this document for warranting a passing grade.

References

1. stat(1): man 1 stat
<http://linux.die.net/man/1/stat>
2. stat(2): man 2 stat, or man fstat
<http://linux.die.net/man/2/stat>
3. md5sum(1): man 1 md5sum
<http://linux.die.net/man/1/md5sum>
4. evp_digestinit(3): man 3 evp_digestinit
http://linux.die.net/man/3/evp_digestinit
5. sha1sum(1): man 1 sha1sum
<http://linux.die.net/man/1/sha1sum>
6. sha1(3): man 3 sha1
<http://linux.die.net/man/3/sha1>
7. Bourne Shell Programming, Steve Parker
<http://steve-parker.org/sh/sh.shtml>

8. Bourne Shell Programming, Andrew Arensburger
<http://www.ooblick.com/text/sh>
9. UBUNTU Hypertext Man Pages
<http://manpages.ubuntu.com/>
10. Linux Shell Scripting Tutorial: A Beginner's handbook
<http://www.freeos.com/guides/lsst/>
11. The Linux Documentation Project
<http://www.tldp.org/>
12. Advanced Bash-Scripting Guide
<http://tldp.org/LDP/abs/html/>
13. A collection of tutorials
<http://www.grymoire.com/Unix/>
14. The Bash Manual
<http://www.gnu.org/software/bash/manual/bash.html>